

Tier 1 Model Inventory Record

HarperCollins Manuscript Evaluator (HC-QG-001)

Document Status: Approved for Pilot

Version: 4.0

Date: April, 2026

Business Owner: Editorial Director

Model Owner: Manager, Editorial Systems

Second Line (Oversight): Model Risk Management (MRM) Team

Section 1: General Information

Model Name: HC Manuscript Evaluator v1

Model ID: HC-QG-001

Model Status: Approved for Pilot (pending final signatures)

Intended Use Case: To support (not replace) acquisitions editors in evaluating manuscript query letters by providing structured, AI-generated analysis.

Section 2: Model Description & Scope

Model Type:

Generative AI (LLM) within a human-in-the-loop decision system.

System Description:

An editor submits a query letter via a controlled prompt template to the OpenAI API. The system returns structured feedback (score, comps, strengths/weaknesses). The editor must review all outputs, complete a mandatory verification checklist, and make the final, independent decision.

Key Design Principle: AI is advisory only; all decisions remain human-controlled and accountable.

Access Method (Production):

- API Only: with audit logging, zero data retention configuration, and pinned model version.
- Web interface explicitly prohibited for production use.

Inputs:

- Query letter text (includes author name, title, genre, summary).
- Standardized prompt template

Outputs:

- Marketability score (1–10)
- Comparable titles
- Strengths and weaknesses
- Recommendation

Dependencies:

- External: OpenAI API (GPT-4 class, pinned version).
- Internal: Human Review Checklist, Audit Logging System, Prompt Template Library.

Section 3: Risk Assessment & Tiering (MRM Approved)

Risk Factor Scoring (Inherent Risk Assessment):

Factor	Risk Rating (0–2)	Rationale
Decision Criticality	2	The model's output materially influences a high-stakes business decision (manuscript acquisition) with significant financial and career impact. The human reviewer is a control, not a factor in assessing inherent risk.
Regulatory Materiality	2	The process is subject to broad regulations including copyright law and anti-discrimination laws. Model failures (e.g., bias, plagiarism) could lead to direct regulatory and legal violations, creating high inherent risk.
Data Sensitivity	2	Processes unpublished manuscripts and author-identifying information (confidential IP and potential PII). This is among the most sensitive data the company owns.
Model Complexity	2	LLM is a non-deterministic, "black-box" system with limited explainability, making it inherently difficult to predict or fully control.
Customer Impact	2	Outputs directly affect authors' opportunities and publishing outcomes. A biased or incorrect evaluation

		has a direct, high impact on a key stakeholder group.
--	--	---

TOTAL 10/10

Tier Classification:

Tier 1 (High Risk)

Justification:

The model's perfect inherent risk score of 10/10 on the framework mandates a Tier 1 classification. This is driven by its material influence on critical decisions, use of sensitive data, and exposure to significant legal and reputational risks. The highest level of governance is required.

Condition for Use:

This model is approved for pilot on the condition that all Tier 1 controls, as documented in this inventory and supporting artifacts, are implemented and remain effective. Failure to meet the quantitative monitoring thresholds defined in Section 8 will result in the immediate suspension of the model's use, pending a joint review by the Model Owner and MRM.

Section 4: Key Risk Themes

The primary risk categories governing this model are:

- Bias Risk: Outputs may vary based on author-identifying information.
- Hallucination Risk: Fabricated comparable titles or analysis.
- Copyright Risk: Potential reproduction of protected text.
- Automation Bias Risk: Editors over-trusting model outputs.
- Auditability Risk: Lack of traceability without proper logging.
- These risks are formally controlled and tested in separate governance artifacts (Risk Control Matrix, Validation Pack, Monitoring Specification).

Section 5: LLM-Specific Metadata

Field	Entry
Base LLM	OpenAI API (GPT-4 class, pinned version)
Access Method	API (audit logging enabled, data retention off)
Prompting Strategy	Standardized template (zero-shot)
Prompt Management	Centralized template library (version controlled)

Section 6: Control Environment (Summary)

As a Tier 1 model, this system relies on a layered control framework:

Preventive Controls:

- Human review checklist (mandatory per use).
- Standardized prompt templates.
- Pre-flight prompt instruction: "Do not reproduce manuscript text verbatim. Summarize only."

Detective Controls:

- Bias monitoring via monthly redaction tests.
- Hallucination checks via random sampling of comparable titles.
- Checklist completion rate monitoring.

Corrective Controls:

- Human override of outputs (must be documented with reasoning).
- Escalation through editorial/legal review for outputs flagged as high risk.

Section 7: Data Considerations

Data Sensitivity:

- Unpublished manuscripts (Confidential IP).
- Author-identifying information (PII).

Production Controls:

- API with data retention disabled by vendor agreement.
- Audit logging access restricted to MRM and authorized personnel.
- All data encrypted in transit (TLS 1.2+).

Section 8: Monitoring & Validation Linkage

Quantitative Performance Anchors (Tier 1 Requirements):

Metric	Acceptable	Threshold
Hallucination rate (comps)	<5% fabricated	Monthly
Bias effect (score change)	≤0.5 point avg. difference	Monthly
Human override rate	30%–70% acceptable range; breach for 2 consecutive months triggers review	Quarterly
Checklist completion rate	100% of submissions	Monthly

Audit log integrity	>99.9% complete	Quarterly
Copyright Flag rate	0 instances	Per incident

Note: The Human Override Rate range is designed to detect both automation bias (low override rates) and model misalignment (high override rates). The Copyright Flag Rate threshold is zero; any single instance triggers an immediate review by Legal Compliance.

Linked Governance Artifacts:

This inventory serves as the central index linking to:

- Monitoring Specification: Defines thresholds and escalation procedures.
- Validation Evidence Pack: Contains results from pre-pilot bias, hallucination, and copyright testing.
- Audit Logs: API-based, immutable record of all transactions.

Section 9: Human Oversight

- All outputs must be reviewed using the mandatory checklist before use.
- Editors retain full decision authority and accountability.
- Model outputs cannot be used as the sole justification for any decision.
- Monthly calibration sessions between editors and MRM to review override patterns and monitoring results.

Section 10: Approvals

Role	Signature	Date
Model Governance Lead	_____	April 2026
Model Risk Management (MRM) Lead	_____	April 2026
Editorial Director (Business Owner)	_____	April 2026

Final Assessment

The HC Manuscript Evaluator is a Tier 1 (High Risk) AI system approved for pilot use, requiring strong governance, independent validation, and continuous monitoring. This status is conditional on ongoing compliance with the controls and monitoring thresholds defined herein. Any material change to the model's scope, access method, or prompting strategy requires a formal review by MRM and an update to this record.